

Research Report

2026 S.T. Yau High School Science Award (Asia)

Bundle Neural Networks for Functional Connectomes

Density-Dependent Aggregation and Cross-Site Prediction

Team Member	Adit Patil
School	School of Scholars
City and Country	Akola, India
Registration Number	Comp-183
Supervising Teacher	Anjali Tupe
Position	Class Teacher
Supervising School	School of Scholars, Akola, India
Report Version	Final evidence report with E1 and E2 mechanism studies, 14 August 2026

Completed evidence report. This version reports the frozen ABIDE-I operator study, the audited post-hoc checkpoint intervention analysis, and a separately labelled synthetic known-geometry study. External validation and additional architecture studies are not presented as completed results.

This document is a separate report-format rendering. It does not replace the existing IEEE-style manuscript or the accepted evidence snapshots. Before an external submission, the researcher must verify the personal details, complete the acknowledgments and AI-use disclosure, and sign the integrity commitments.

Abstract

Functional connectomes have a natural graph representation, but the usual node feature is already a complete connectivity profile. Neighborhood aggregation may therefore blur useful differences instead of supplying new information. This study tested whether learned orthogonal bundle transport changes that density-dependent behavior. The analysis used 754 technically eligible ABIDE-I participants from 18 sites. A shared neural backbone was paired with identity propagation, a learned-local capacity control, Graph Convolutional Network (GCN) aggregation, trivial-bundle diffusion, or learned Bundle Neural Network (BuNN) transport. Nonzero graphs retained the strongest positive edges at 1, 5, 10, and 20 percent density. Evaluation used nested held-out-site validation, five final seeds, regularized non-graph baselines, and common-frame representation diagnostics.

The primary BuNN-minus-GCN performance-curve contrast was -0.00958 (95% site-bootstrap confidence interval $[-0.03665, 0.01146]$; exact sign-flip $p = 0.524$). Relative to the connectome elastic net, the BuNN contrast was -0.05516 (95% confidence interval $[-0.08297, -0.02752]$). The matched-anchor effective-rank contrast was -0.00719 (95% confidence interval $[-0.01310, -0.00105]$), which ran against the proposed representation-preservation advantage. Performance declined as graph density increased for each diffusion operator. Site weighting and training seed also affected the small BuNN–GCN difference.

A separately frozen post-hoc intervention analysis then tested what the 360 accepted BuNN checkpoints actually used. Replacing learned maps with random orthogonal maps reduced equal-site balanced accuracy by 8.86 percentage points (95% interval $[-12.36, -4.78]$; Holm-adjusted exact $p = 0.0038$). Identity maps, shuffled learned maps, and degree-preserving graph rewiring caused smaller reductions and did not pass the four-contrast multiplicity-controlled test. Representation changes were substantial but only weakly associated with prediction changes across site-density cells.

A second post-hoc study tested the proposed mechanism where the true topology and coordinate frames were known. Across 700 frozen synthetic cells, learned BuNN gained 42.75 percentage points on the pre-specified conditional contrast (95% paired bootstrap interval $[38.42, 47.67]$; exact $p = 0.00195$). The oracle, matched learned-local, no-geometry, wrong-topology, shuffled-marker, unlearnable-frame, and global-feature controls all behaved according to their pre-specified roles. Learned BuNN solved the recoverable-geometry task but did not exactly recover the planted relative transports.

Under this specified ABIDE-I pipeline, learned bundle transport provided no detected advan-

tage over GCN or the elastic net. The intervention audit shows that its learned maps nevertheless influenced its own predictions; the exact graph topology contributed much less. The study does not establish a biological bundle geometry, a clinical diagnostic system, or a universal result for BuNNs. Its contribution is a controlled operator comparison and a checkpoint-level audit in a dense, globally informed connectome setting, followed by a controlled demonstration of when transport can and cannot help.

Keywords: functional connectome; bundle neural network; graph neural network; representation collapse; held-out-site validation; ABIDE-I; autism

Acknowledgments and Assistance Statement

The study used public, anonymized ABIDE-I derivatives distributed through the Preprocessed Connectomes Project. The original ABIDE investigators and participating institutions collected and shared those data; they were not involved in the present analysis. The computational work used open-source Python, PyTorch, NumPy, pandas, scikit-learn, and scientific plotting tools, with training performed on an AWS EC2 instance equipped with an NVIDIA A10G GPU.

AI systems assisted with literature discovery, protocol criticism, code development and debugging, run monitoring, analysis checks, figure and table production, and document formatting. The project maintains a dated local record in `docs/ai_use_log.md`. AI assistance does not establish authorship or responsibility for the scientific claims. Before external submission, the student researcher must verify the analysis, rewrite or edit the report in their own voice, and complete any missing tool versions, dates, frequencies, supervisor approvals, and retained chat records required by the competition’s AI-use policy.

System	Recorded assistance	Status before submission
OpenAI Codex	Research planning, literature leads, implementation and test assistance, AWS run control, analysis auditing, plots, and report formatting.	Detailed dated entries exist locally; exact product/model version and supervisor approval must be verified.
Anthropic Claude	Methodological critique and proposal discussion reported by the researcher.	Exact version, dates, prompts, and retained records must be supplied by the researcher.
Cursor	Interactive coding assistance reported by the researcher.	Exact models, dates, tasks, and retained records must be supplied by the researcher.

Supervisor contribution. The researcher must replace this sentence with an exact account of the supervising teacher’s intellectual, technical, editorial, and logistical contributions. No contribution should be inferred or invented.

Research Integrity Commitments

The following page is included for format readiness. It is not signed in this completed evidence draft.

1. The reported data, code outputs, analyses, and citations must be honest, traceable, and presented without selective omission.
2. The student researcher must understand and be able to explain the methods, implementation, results, limitations, and conclusions.
3. All external assistance, including AI assistance, must be disclosed in accordance with the applicable competition rules.
4. Public data and software must be credited, and participant-level data must not be exposed in the report or release package.
5. No incomplete, exploratory, or planned experiment may be described as a completed result.
6. Any relationship that could create a conflict of interest with judging or supervision must be disclosed.
7. The work must comply with relevant safety, research-integrity, and competition requirements.
8. The student and supervisor accept responsibility for checking the final submission rather than relying on automated tools.

Student signature: _____

Date: _____

Supervisor signature: _____

Date: _____

Contents

Abstract	i
Acknowledgments and Assistance Statement	iii
Research Integrity Commitments	iv
1 Introduction	1
1.1 How the original question was refined	1
2 Background	2
2.1 ABIDE-I and cross-site prediction	2
2.2 Graph convolution and aggregation	2
2.3 Bundle transport and coordinate frames	2
2.4 Why published ABIDE percentages are not one leaderboard	3
3 Methods	3
3.1 Study design and frozen evidence boundary	3
3.2 Dataset and technical quality control	4
3.3 Connectome construction	4
3.4 Classical baselines	6
3.5 Shared neural backbone and operator controls	6
3.6 Nested held-out-site evaluation	7
3.7 Density-curve and representation estimands	7
3.8 Statistical analysis and decision rule	8
3.9 Run integrity and reproducibility	8
3.10 Post-hoc accepted-checkpoint intervention audit	8
3.11 Synthetic known-geometry mechanism study	9
4 Results	10
4.1 Classical baseline performance	10
4.2 Predictive performance across graph density	10
4.3 Common-frame representation behavior	12
4.4 Accepted-checkpoint mechanism audit	13

4.5	Synthetic known-geometry results	15
4.6	Site, seed, and weighting sensitivity	18
4.7	Execution cost	18
5	Discussion	20
5.1	Why the result was plausible before the experiment	20
5.2	What the completed evidence can explain	21
5.3	Why the learned-local control matters	21
5.4	Site heterogeneity and uncertainty	21
5.5	Limitations	23
5.6	Remaining extensions and their evidential role	24
6	Conclusion	24
	Evidence Boundary and Remaining Work	26
A	Reproducibility Record	27
B	AI-Use Detail Required Before Submission	27

1 Introduction

Resting-state functional magnetic resonance imaging records slow changes in the blood-oxygen-level-dependent signal while a participant is not performing a specific task. A functional connectome summarizes the statistical association between pairs of regional time series. Brain regions become graph nodes, and correlations between their signals become weighted edges. This representation makes graph neural networks an attractive modeling choice.

That attraction contains a hidden assumption: message passing must add useful information. In many connectome classifiers, each node already receives one full row of the correlation matrix. Its feature vector therefore describes its association with every region in the atlas. Aggregating neighboring rows does not introduce context that was previously absent. It may instead average away distinctions that a classifier needs.

Recent controlled benchmarking found that classical models and non-graph neural networks can match or outperform graph deep-learning models on several functional-connectome tasks. Ablations in that work implicated message aggregation as one source of degradation [1]. A graph representation alone therefore does not justify graph propagation. The operator must be compared with no propagation and with strong regularized baselines on the same data splits.

BuNNs offer a distinct propagation rule. Before diffusion, they use learned orthogonal maps to transport node features into compatible coordinate frames [2]. Theoretical and synthetic results suggest that this construction can resist representation collapse more effectively than ordinary diffusion. Neural-sheaf methods provide a related example of how non-trivial transport can alter message passing [3]. The open question is whether this advantage carries into dense functional connectomes with globally informed node features.

1.1 How the original question was refined

The first formulation of this project treated negative functional correlations as a possible form of heterophilic message passing. Formalizing the experiment showed that this interpretation was not supportable. Correlation sign does not identify graph-learning heterophily, anatomical connections, neural excitation or inhibition, or causal flow. It can also change with preprocessing choices such as global-signal regression [4]. The original dataset, connectome representation, BuNN architecture, and aggregation problem were retained, but the hypothesis was narrowed to the computational behavior of the propagation operator.

The resulting study asks:

On a pre-specified ABIDE-I pipeline, does learned bundle transport produce a more favorable held-out-site performance–density curve than GCN aggregation, and does any predictive pattern co-occur with stronger common-frame representation preservation?

An advantage was required to exceed GCN, matched no-propagation controls, and the strongest regularized connectome baseline. This stricter rule prevents a complex operator from being credited merely because it is less harmful than another graph operator.

2 Background

2.1 ABIDE-I and cross-site prediction

ABIDE-I combined resting-state fMRI and phenotypic data from international sites to support large-scale autism research [5]. The Preprocessed Connectomes Project later distributed derivatives produced by several pipelines, including C-PAC [6]. Multi-site prediction is difficult because scanners, acquisition protocols, participant composition, motion, and preprocessing can differ across sites. A random participant split can place nearly identical site conditions in training and test data. Holding out an entire site creates a stricter evaluation: the model must generalize to a data-collection environment it did not see during fitting.

2.2 Graph convolution and aggregation

A GCN combines a node’s features with normalized features from neighboring nodes and then applies a learned transformation [7]. As propagation repeats or the graph becomes denser, representations can contract. Whether this is harmful depends on the graph, input features, architecture, and task. Here the question is especially relevant because each node begins with a complete connectivity row. A no-propagation control is therefore part of the scientific comparison, not merely an ablation.

2.3 Bundle transport and coordinate frames

BuNNs associate node features with local coordinate frames. Learned orthogonal maps align features before diffusion [2]. This creates an important measurement issue. Raw embeddings at different nodes may look dissimilar simply because they are expressed in different frames. Directly stacking them and computing cosine similarity or effective rank could reward rotations instead of retained task information. The present study maps embeddings into a common learned

frame before comparing nodes and also uses a learned-local control that retains map-generating capacity without inter-node diffusion.

2.4 Why published ABIDE percentages are not one leaderboard

Published graph-model results on ABIDE answer different questions. Parisot et al. reported 70.4% accuracy and 0.75 AUC on 871 participants, but their graph connected participants rather than brain regions. Its edges used phenotypic information including sex and acquisition site, and evaluation used 10-fold transductive cross-validation [8]. Wang et al. reported 71.6% mean leave-one-site-out accuracy on 1,057 participants. Their cGCN used CCS-preprocessed CC200 ROI time series, a training-derived k -nearest-neighbour graph, five graph-convolution layers, and temporal averaging [9]. Yang et al. reported 64.27% average accuracy for a graphlet-augmented GCN on 1,035 participants, together with a 75.9% highest site-specific value [10].

None of those percentages is numerically interchangeable with the present equal-site balanced accuracy. Cohort filtering, atlas, preprocessing, input representation, graph definition, split, metric, and model selection all differ. The released Wang leave-one-site-out script also passes the left-out site to Keras as validation data and uses its validation performance for stopping and checkpoint selection. This code-level leakage does not erase the published result, but it prevents 71.6% from serving as a clean unseen-site target for this study. Han et al. provide the closest conceptual comparison: their controlled benchmark found that simple models often matched or exceeded graph models and that aggregation performance declined with density [1].

3 Methods

3.1 Study design and frozen evidence boundary

The experiment isolated the propagation operator while holding the rest of the pipeline fixed. Classical baselines were completed and audited first. Five neural variants then used the same inputs, outer test sites, grouped inner folds, candidate budget, stopping rule, and final seeds. Before results were opened, the complete neural artifact was checked for expected coverage and the confirmatory analysis was frozen. The accepted run contained 9,324 neural fits, 52,780 held-out predictions, and results for all 18 outer sites.

Study 1 is the confirmatory ABIDE-I comparison. This report also includes two separately frozen post-hoc studies: E1 intervenes on the 360 accepted BuNN checkpoints, and E2 tests the proposed computation on synthetic graphs with known topology and coordinate frames. Their

Table 1: Protocol context for selected ABIDE graph-model results. Reported percentages are retained in their original metrics and should not be ranked as if they came from one experiment.

Study	Data and representation	Evaluation and result	Main mismatch with this study
Present study	754; C-PAC no-GSR; AAL-116 connectivity-row brain graphs	Nested leave-one-site-out; equal-site BA; elastic net 0.6401, GCN curve 0.5945, BuNN curve 0.5849	Primary outcome is a balanced-accuracy density curve with test sites excluded from model selection.
Wang et al. [9]	1,057; CCS; CC200 raw ROI time series; training-derived k -NN graph	Leave-one-site-out ordinary accuracy; 71.6% at $k = 5$	Different input, atlas, architecture, cohort, and metric; released code uses the left-out site for model selection.
Parisot et al. [8]	871; participant population graph with imaging and phenotypic edges	10-fold transductive CV; 70.4% accuracy, 0.75 AUC	Nodes are participants; graph uses site and sex; folds mix sites.
Yang et al. [10]	1,035; functional graphs and graphlet counts	64.27% average accuracy; 75.9% highest site value	Different features and aggregation; a highest-site value is not an overall estimate.
Han et al. [1]	1,009; C-PAC; CC200 connectomes	AUROC benchmark and density ablations	Best used as a relative aggregation comparison, not an absolute score target.

protocols were fixed only after Study 1 was known, so they explain the completed result but do not independently confirm a mechanism in ABIDE-I. Deeper networks, Wang-style time-series training, and ABIDE-II evaluation remain outside the completed findings.

3.2 Dataset and technical quality control

The source was the ABIDE-I C-PAC band-pass-filtered AAL ROI time-series derivative without global-signal regression. The AAL atlas divides the brain into 116 macroscopic regions in standardized space [11]. Eligibility required a valid file header, finite numeric values, nonzero variance in every ROI time series, and a finite resulting connectome. Fifteen of 769 downloaded files contained at least one zero-variance ROI, making the correlation undefined. Their exclusion left 754 technically eligible participants: 371 labeled ASD and 383 controls across 18 sites.

Technical eligibility does not imply population representativeness. Site composition remained heterogeneous, as shown in figure 1. That variation is handled in evaluation and uncertainty estimation rather than erased from the data.

3.3 Connectome construction

For each participant, Pearson correlations were computed between every pair of ROI time series and transformed with Fisher’s r -to- z function. The diagonal was set to zero. The resulting

Table 2: Frozen cohort and feature summary.

Quantity	Value
Downloaded ABIDE-I ROI time-series files	769
Zero-variance-ROI technical exclusions	15
Technically eligible participants	754
ASD / control participants	371 / 383
Evaluable acquisition sites	18
AAL regions per participant	116
Unique undirected connectome features	6,670




Figure 1: Participant composition of the frozen technically eligible cohort. Sites varied in both total size and class balance, motivating site-level evaluation and equal-site aggregation. This is a descriptive count figure, not a performance result.

matrix was finite and symmetric. Its 6,670 lower-triangle entries formed the feature vector for classical models. For neural models, each node received its full 116-value connectivity row.




Figure 2: A deterministic example from the frozen archive illustrating a participant-level Fisher- z functional-connectivity matrix. Red and blue indicate positive and negative statistical associations. They are not interpreted as excitation, inhibition, anatomy, or causal flow.

At nonzero density, graphs retained the strongest positive off-diagonal associations symmetrically at 1, 5, 10, and 20 percent. Density zero supplied the matched identity or learned-local anchor. Negative associations remained part of the node connectivity profiles but were not used as propagation edges. This was a pre-specified computational choice, not a biological claim.

3.4 Classical baselines


Three logistic models provided practical references. The first used age, sex, mean framewise displacement, and usable scan length. The second used all vectorized connectome features with elastic-net regularization, which combines ℓ_1 and ℓ_2 penalties for correlated high-dimensional predictors [12]. The third combined connectome features and covariates. All scaling, feature processing, and tuning were fitted on training sites only.

3.5 Shared neural backbone and operator controls

The neural variants shared the encoder, hidden width, two propagation layers, nonlinearities, normalization, dropout, pooling, classifier, optimizer, training schedule, and tuning budget. The controlled difference was the propagation rule or the map capacity needed to isolate it.

Table 3: Neural operator and capacity controls.

Variant	Operation	Purpose
Identity	No inter-node propagation	Tests whether message passing adds value.
Learned-local	Learns the BuNN node-wise maps but does not diffuse between nodes	Separates map-generator capacity from transport.
GCN	Normalized neighborhood aggregation	Standard graph-diffusion reference.
Trivial bundle	Bundle diffusion with identity transport maps	Separates the diffusion construction from learned transport.
Learned BuNN	Learned orthogonal node maps followed by bundle diffusion	Tests the proposed transport operator.

**Figure 3:** Computational distinction among GCN aggregation, the learned-local capacity control, and learned bundle transport. The schematic does not represent biological signal flow.

3.6 Nested held-out-site evaluation

Each site served once as the outer test site. Four grouped folds within the remaining sites selected among four equal-budget AdamW candidates using two tuning seeds. The selected configuration was refitted with five fixed final seeds. Test-site data did not enter scaling, tuning, stopping, or model selection. The primary predictive endpoint was the unweighted mean of held-out-site balanced accuracies. This gives each acquisition site equal influence. Pooled balanced accuracy and AUROC were secondary.

3.7 Density-curve and representation estimands

The primary predictive estimand was the normalized trapezoidal area under the balanced-accuracy curve across the full density grid. GCN and trivial-bundle curves used identity at zero density. The learned-BuNN curve used the parameter-matched learned-local anchor. Comparing the whole curve avoided choosing a favorable density after viewing the results.

After each propagation layer, BuNN embeddings were transported into a common learned frame before node comparisons. Normalized effective rank at layer 2 was the primary representation endpoint. Secondary measures included normalized dispersion, mean pairwise cosine similarity, invariant edge-transport distance, and encoder-to-layer-2 rank change. These measures describe co-occurring representation behavior. They do not prove that a representation change caused a prediction change.

3.8 Statistical analysis and decision rule

The five final seeds were averaged within each site and configuration before site-level inference. Confidence intervals used 10,000 paired bootstrap resamples of the 18 held-out sites. Curve contrasts also received exact two-sided paired sign-flip tests. Density-specific tests were secondary and Holm-corrected within each named family.

A positive architectural result required all three of the following:

1. stronger common-frame representation preservation than GCN;
2. better held-out prediction than GCN and the matched no-propagation controls; and
3. performance above the strongest regularized non-graph baseline.

A confidence interval containing zero was reported as no detected advantage. A practically meaningful gain was ruled out only if the upper confidence bound fell below the pre-specified 0.03 balanced-accuracy margin.

3.9 Run integrity and reproducibility

Long-running jobs wrote atomic per-site checkpoints and recorded immutable hashes for data, configuration, and source code. Site ownership, restart validation, canonical merging, and independent completion audits protected the parallel run. The neural artifact was checked for coverage, recomputed metrics, diagnostics, runtime records, and warnings before unblinding. Values in the report are generated from a hash-bound result snapshot rather than copied manually.

3.10 Post-hoc accepted-checkpoint intervention audit

After Study 1 was frozen, a separate protocol tested what the 360 accepted learned-BuNN checkpoints depended on without refitting them. Each checkpoint was first required to reproduce its archived held-out probabilities within 10^{-6} . Four scientific interventions were then compared

with unaltered inference: replacing learned maps with identity maps, shuffling learned maps between nodes, replacing them with random orthogonal maps, and rewiring each participant graph by double-edge swaps that preserved its exact degree sequence. An encoded-node permutation served only as an equivariance check.

The randomized families used 100 stored randomizations, shared across the five model seeds. Balanced accuracy was calculated for each seed and randomization, randomizations were averaged within the site-density cell, and seeds were then averaged. The four density-specific paired effects were averaged within site, followed by an unweighted mean over the 18 sites. The effect direction was always intervention minus unaltered BuNN. Uncertainty used 10,000 paired site-bootstrap samples. Two-sided exact sign-flip tests enumerated all 2^{18} site sign assignments, with Holm correction across the four primary interventions.

Secondary endpoints were absolute probability change, classification flips, AUROC, sensitivity, specificity, and changes in the four final-layer gauge-aware diagnostics. Associations between representation and prediction changes were descriptive Spearman correlations across the 72 site-density cells, not mediation tests. The complete run was archived and audited before the analysis protocol and code were frozen. The analyzer passed synthetic known-answer tests before a single acknowledged unblinding invocation, and an independent recomputation audit passed before interpretation.

3.11 Synthetic known-geometry mechanism study

E2 was frozen after Study 1 and E1 and is therefore explanatory rather than confirmatory evidence about ABIDE-I. Each replicate contained 600 balanced synthetic graphs with 116 nodes and fixed 360/120/120 train, validation, and test partitions. Class labels did not change the marginal type of node values: class 1 signals were smooth on a degree-six ring graph, while class 0 used a participant-specific permutation of the same type of smooth signal.

Seven operators shared a fixed identity encoder, one 32-dimensional propagation layer, GELU activation, mean pooling, optimizer, stopping rule, and test protocol: identity, GCN, trivial-bundle diffusion, fixed-random transport, learned-local, learned BuNN, and an oracle supplied with the true maps. The fixed encoder kept the planted coordinate frames identifiable. Learned-local retained BuNN’s map-generator capacity without cross-node diffusion.

The main families were no geometry (S0), fixed recoverable node frames (S1), incorrect topology (S2), shuffled frame markers (S3), subject-specific frames without an informative marker (S5), and a global-feature analogue (S6). Marker-frame corruption was also swept through 15, 30, 60, and 120 degrees. Ten paired replicate seeds produced 700 cells. Test predictions were

opened only after all cells passed structural and hash audits and 100 condition-replicate groups were independently sealed.

The primary estimand was

$$[(\text{BuNN} - \text{GCN})_{S1} - (\text{BuNN} - \text{GCN})_{S0}],$$

measured in test balanced-accuracy percentage points. Support required three ordered gates: a positive oracle-minus-GCN contrast in S1; a primary mean of at least three points with a positive paired-bootstrap lower bound and exact sign-flip $p < 0.05$; and a positive BuNN-minus-learned-local lower bound in S1. A greater-than-three-point BuNN-minus-GCN S0 result triggered a capacity warning. The complete analysis was invoked once and independently recomputed.

The first full launcher was stopped before unblinding after a regression review found that half of its nominal S0 maps were reflections rather than identity maps. Version 2 used explicit identity matrices, added an exact regression test, and restarted every cell without reusing a v1 checkpoint or prediction.

4 Results

4.1 Classical baseline performance

The connectome elastic net was the strongest classical model, with equal-site balanced accuracy 0.6401. Adding covariates produced 0.6336, while covariates alone produced 0.5652. Imaging features therefore carried cross-site predictive information under this pipeline, but the result is computational and should not be described as clinical diagnostic performance.

Table 4: Held-out-site performance of the classical baselines.

Model	Equal-site BA	Pooled BA	Pooled AUROC
Covariates-only logistic regression	0.5652	0.5629	0.5973
Connectome elastic net	0.6401	0.6495	0.6794
Connectome + covariates elastic net	0.6336	0.6428	0.6773

4.2 Predictive performance across graph density

All diffusion curves declined as graph density increased (figure 4). The learned-BuNN minus GCN normalized curve-area difference was -0.00958 (95% confidence interval [-0.03665, 0.01146];

exact $p = 0.524$). The interval crossed zero, so no predictive advantage was detected. Its upper bound was also below the pre-specified 0.03 margin. This rules out a gain of that size for the specified estimand and pipeline, but not every smaller effect or every other BuNN design.




Figure 4: Equal-site mean held-out balanced accuracy across the pre-specified positive-edge density grid. Error bars show 1.96 standard errors across sites for visualization. Confirmatory inference used paired site bootstrap and exact sign-flip tests on the locked curve estimand.

Against the site-matched connectome elastic net, the learned-BuNN curve difference was -0.05516 (95% confidence interval $[-0.08297, -0.02752]$; exact $p = 0.0018$). The interval lay entirely below zero. Comparisons with identity and learned-local controls also failed the positive-result rule.

Table 5: Pre-specified predictive contrasts. Estimates are balanced-accuracy differences.

Contrast	Estimate	95% CI	Exact p
BuNN curve – GCN curve	-0.0096	$[-0.0366, 0.0115]$	0.524
BuNN curve – trivial-bundle curve	-0.0062	$[-0.0293, 0.0108]$	0.681
BuNN nonzero mean – learned-local	-0.0167	$[-0.0347, 0.0009]$	0.096
BuNN nonzero mean – identity	-0.0146	$[-0.0320, 0.0018]$	0.116
BuNN curve – connectome elastic net	-0.0552	$[-0.0830, -0.0275]$	0.0018

4.3 Common-frame representation behavior

The raw BuNN effective-rank curve appeared higher than the GCN curve, but its learned-local zero-density anchor was already high. After subtracting each operator's matched anchor, the normalized effective-rank contrast was -0.00719 (95% confidence interval [-0.01310, -0.00105]). Its direction was opposite to the proposed BuNN preservation advantage. Normalized dispersion and pairwise cosine similarity showed the same broad pattern.




Figure 5: Equal-site common-frame representation diagnostics at layer 2. Zero-density points use the matched identity or learned-local anchor. Raw, unaligned BuNN node embeddings were not compared.

The learned-local control is decisive for interpretation. Without it, extra map-generator capacity could be mistaken for information preserved by inter-node transport. The matched result gives no support for that transport explanation. These diagnostics remain computational measurements and do not identify a biological mechanism.

Table 6: BuNN-minus-GCN matched-anchor representation contrasts.

Layer-2 endpoint	Estimate	95% CI
Normalized effective rank (primary)	-0.0072	[-0.0131, -0.0010]
Normalized dispersion	-0.0417	[-0.0523, -0.0306]
Mean pairwise cosine similarity	0.0415	[0.0319, 0.0507]
Invariant edge-transport distance	-0.6436	[-0.8205, -0.4720]
Encoder-to-layer-2 effective-rank change	-0.0056	[-0.0144, 0.0048]

4.4 Accepted-checkpoint mechanism audit

The post-hoc intervention audit showed that the learned maps were not dormant components. Replacing them with random orthogonal maps reduced held-out balanced accuracy by 8.86 percentage points (95% paired site-bootstrap interval [-12.36, -4.78]; exact sign-flip $p = 0.00095$; Holm-adjusted $p = 0.0038$). The effect was negative at every density and in 15 of 18 sites after averaging densities. This was the only intervention that passed the pre-specified four-contrast multiplicity-controlled test.

Identity maps produced a 4.72-point reduction (95% interval [-8.56, -0.29]). Its unadjusted exact p was 0.0468, but its adjusted p was 0.1404. Shuffling learned maps between nodes caused a smaller 1.94-point reduction (95% interval [-3.77, 0.05]; adjusted $p = 0.1404$). Degree-preserving rewiring caused a 1.23-point reduction (95% interval [-2.49, -0.05]; adjusted $p = 0.1404$). Bootstrap intervals and sign-flip tests answer related but different questions; the adjusted sign-flip result controls the four named primary tests.

Table 7: Accepted-checkpoint intervention effects. Estimates are intervened minus unaltered BuNN balanced accuracy in percentage points.

Intervention	Estimate	95% CI	Holm p
Identity maps	-4.72	[-8.56, -0.29]	0.1404
Shuffled learned maps	-1.94	[-3.77, 0.05]	0.1404
Random orthogonal maps	-8.86	[-12.36, -4.78]	0.0038
Degree-preserving rewiring	-1.23	[-2.49, -0.05]	0.1404

The random-map effect remained negative at 1, 5, 10, and 20 percent density (figure 7). Identity maps had their largest effect at 1 percent density. Shuffling maps and rewiring topology produced smaller effects that often had density-specific intervals crossing zero. The density pattern therefore does not look like a benefit confined to one hand-picked threshold.

The prediction changes were also practically visible. Random maps changed 49.46% of binary decisions and reduced AUROC by 14.59 points. Identity maps changed 34.41% of decisions and




Figure 6: Overall E1 intervention effects with 95% paired site-bootstrap intervals. Negative values mean that disrupting the component reduced the accepted checkpoint’s held-out balanced accuracy. Statistical decisions use the Holm-adjusted exact sign-flip tests in table 7.




Figure 7: Density-specific accepted-checkpoint intervention effects. Shaded regions are 95% paired site-bootstrap intervals. These density-specific estimates are descriptive; the primary inference averages the four effects within site before equal-site analysis.

reduced AUROC by 9.50 points. Shuffled maps changed 15.12% of decisions, while topology rewiring changed 6.96%. Random maps strongly reduced sensitivity while increasing specificity, indicating a shift toward control predictions rather than symmetric noise.

All four interventions changed at least some gauge-aware diagnostics, but those changes did not track predictive effects closely. Across the 72 site-density cells, the absolute descriptive Spearman correlations between diagnostic changes and balanced-accuracy changes were at most 0.21 (figure 8). The measured representation summaries therefore do not provide a simple predictive mechanism or a causal mediation result.




Figure 8: Descriptive association between final-layer gauge-aware diagnostic changes and balanced-accuracy changes across 72 site-density cells. The correlations are exploratory summaries, not multiplicity-controlled tests or evidence of causal mediation.


4.5 Synthetic known-geometry results

The frozen E2 classification was *supported*. The primary conditional learned-BuNN advantage was 42.75 percentage points (95% paired bootstrap interval [38.42, 47.67]; exact sign-flip $p = 0.001953$). This is the increase in the BuNN-minus-GCN difference from S0 to S1, paired across the ten replicate seeds. The point estimate exceeded the three-point practical threshold and its interval excluded zero.

The oracle positive control and matched-capacity gate passed, while the S0 capacity warning did not trigger (table 8 and figure 9). In S0, GCN and learned BuNN both reached approximately perfect performance (1.000 and 0.998), showing that ordinary aggregation solved the task when frames were aligned. In S1, learned BuNN reached 0.990 balanced accuracy and the oracle reached 1.000; GCN reached 0.565 and learned-local reached 0.503.

Table 8: Pre-specified E2 validity gates. Differences are test balanced accuracy in percentage points.

Frozen contrast	Estimate	95% CI	Exact p
Conditional BuNN advantage	42.75	[38.42, 47.67]	0.001953
Oracle – GCN in S1	43.50	[39.42, 48.00]	0.001953
BuNN – learned-local in S1	48.67	[45.83, 51.50]	0.001953
BuNN – GCN in S0	−0.25	[−0.58, 0.00]	0.500000

**Figure 9:** E2 paired validity contrasts with 95% bootstrap intervals across ten replicate seeds. The first three contrasts are ordered support gates; the S0 contrast checks for a generic BuNN advantage when alignment is not required.

The corruption controls restricted that success to the intended boundary conditions (figure 10). Using the wrong topology reduced learned-BuNN balanced accuracy to 0.493 and even reduced the oracle to 0.517. Shuffling the label-independent frame marker reduced learned BuNN to 0.503, while the oracle, which still received the true signal maps, remained at 1.000. When frames changed independently for every participant without an informative marker, learned BuNN reached 0.498 and the oracle reached 1.000. The paired S1 reductions for wrong topology, shuffled markers, and unlearnable frames were 49.67, 48.67, and 49.17 points respectively; all exact $p = 0.001953$.




Figure 10: Mean test balanced accuracy across E2 conditions. S1 provides fixed recoverable geometry; S2, S3, and S5 break topology, marker assignment, and across-participant learnability. S6 supplies an already-global summary.

The marker-noise sweep was graded rather than binary (figure 11). Learned-BuNN mean balanced accuracy was 0.990 at zero degrees, 0.992 at 15 degrees, 0.990 at 30 degrees, 0.891 at 60 degrees, and 0.511 at 120 degrees. The 60-degree result varied substantially across seeds (SD 0.202), so the curve demonstrates a corruption boundary without defining a universal angular tolerance.

S6 supplied a replicated whole-graph statistic to every node. Identity, GCN, trivial bundle, fixed-random transport, learned-local, learned BuNN, and the oracle all reached 1.000 balanced accuracy. This control shows directly that message passing can become redundant when node inputs already disclose the relevant global information.

Learned BuNN did not recover the planted maps exactly. In S1, its mean gauge-invariant relative-transport error was 22.07, compared with 30.62 for identity-frame operators and zero




Figure 11: Performance across corruption of the frame marker. The oracle continues to receive true signal maps; learned BuNN must infer a useful transport from progressively less reliable marker information.

for the oracle. The model learned a task-useful transport solution, but near-perfect classification is not evidence of exact map identification.

4.6 Site, seed, and weighting sensitivity

Only 1 of the 18 leave-one-site influence estimates was positive, and no favorable exclusion interval lay above zero. Two of five seed point estimates favored BuNN, but no favorable seed interval excluded zero; one unfavorable interval lay below zero. Participant weighting reversed the sign of the small BuNN–GCN point estimate, whereas the equal-site mean and median remained negative. The pre-specified confirmatory estimand was the equal-site mean.

4.7 Execution cost

Learned BuNN used 8,529 parameters, compared with 5,889 for GCN. Across the recorded fits, BuNN required 31.14 total fit-hours and GCN required 15.09. Maximum recorded peak GPU memory remained below 0.38 GiB for both. The low memory use reflects small graphs and models; the workload was dominated by many independent fits rather than a single large tensor. Runtime figures are specific to this implementation and the A10G system.




Figure 12: Leave-one-site influence on the BuNN–GCN curve contrast. This analysis tests whether the confirmatory average is driven by a single site; it does not create 18 independent replications.




Figure 13: BuNN–GCN contrast across the five fixed final training seeds. Variation across seeds shows that the small average difference is sensitive to optimization randomness.

Table 9: Observed operator cost and equal-site curve performance.

Operator	Parameters	Fit-hours	Peak GPU GiB	Curve BA
GCN	5,889	15.09	0.373	0.5945
Trivial bundle	5,889	25.99	0.377	0.5911
Learned BuNN	8,529	31.14	0.378	0.5849

5 Discussion

The proposed BuNN advantage did not carry into this ABIDE-I pipeline. Learned bundle transport did not improve the performance–density curve over GCN, did not beat the strongest regularized connectome baseline, and did not preserve matched-anchor effective rank more strongly. Because inputs, backbone, tuning budget, and operator controls were held fixed, the result speaks directly to this propagation design rather than to a broad comparison of unrelated models.

The intervention audit resolves one ambiguity in that negative comparison. BuNN’s learned maps were not ignored: replacing them with random orthogonal maps caused a large, multiplicity-controlled loss of held-out performance. The failure to beat GCN or the elastic net therefore cannot be explained by saying that training never used the map system. Instead, computationally important learned transformations existed inside a model whose overall inductive bias still did not improve cross-site prediction.

E2 shows that this distinction is coherent rather than contradictory. With a correct graph and fixed recoverable coordinate frames, learned BuNN sharply outperformed GCN and its learned-local capacity control. The advantage vanished when topology was wrong, marker geometry was shuffled, or frames varied across participants without recoverable information. Bundle transport can therefore solve the mechanism it was designed for, while still adding no value to the frozen ABIDE-I representation.

5.1 Why the result was plausible before the experiment

The original hypothesis was not arbitrary. Ordinary diffusion can contract node representations, while orthogonal transport can align features before averaging. In graph problems where neighboring features are expressed in different useful frames, that operation can protect structure that direct averaging would erase. BuNN theory and prior synthetic evaluations therefore provided a reasonable architectural motivation [2].

The experiment nevertheless placed that idea in an unusually demanding setting. Functional connectomes are dense, and the node feature is already a global profile. The benefit of transport-

ing locally available information may be small when every node already describes its relationship with all other nodes.

5.2 What the completed evidence can explain

The data support a layered explanation, not a single proven mechanism.

The strongest current explanation is therefore architectural and data-based. The full connectivity rows already expose global information, while denser propagation averages those profiles. Learned orthogonal maps materially shape BuNN’s own predictions, but the much smaller topology-rewiring effect suggests that message passing over the selected positive-edge graph contributes less. The model can learn useful internal coordinate transformations without those transformations outperforming a simpler GCN or a linear model that directly uses the complete connectome. E2 strengthens this interpretation by separating the two missing requirements. Learned BuNN succeeded when the graph actually carried the relevant smoothness relation and frame information was recoverable; it failed when either condition was broken. When a global summary was supplied at every node, every operator succeeded without a propagation advantage. This interpretation fits Study 1, E1, and E2; it is not a causal proof that the same feature components or coordinate frames exist in fMRI.

5.3 Why the learned-local control matters

BuNN introduces additional parameters through its map generator. A comparison with GCN alone could credit transport for differences caused by that added capacity. The learned-local control keeps the map generator but removes inter-node diffusion. Its high zero-density effective rank shows how a naive raw-embedding analysis could overstate preservation. Anchoring each operator to its corresponding no-propagation state makes the mechanism comparison more credible.

5.4 Site heterogeneity and uncertainty

The primary BuNN–GCN confidence interval crossed zero, and participant weighting reversed the point-estimate sign. Individual site exclusions and training seeds changed the magnitude and sometimes the direction of this small difference. The correct conclusion is not that the models are exactly equivalent. It is that no BuNN advantage was detected and that a gain of 0.03 balanced accuracy on the pre-specified curve estimand is inconsistent with the observed upper confidence bound.

Table 10: Evidence-based interpretation of the null result.

Observation	Supported interpretation	What it does not prove
Diffusion curves decline with density	Adding more propagation edges is associated with worse held-out-site prediction in this pipeline.	A universal law of over-smoothing or a causal biological mechanism.
BuNN does not beat GCN	Learned transport provides no detected predictive advantage over ordinary aggregation here.	That BuNN never works on brain or non-brain graphs.
BuNN remains below elastic net	A regularized model on the full connectome is the stronger practical reference for these inputs.	That nonlinear modeling is always unnecessary.
Learned-local anchor already has high raw rank	Extra node-wise map capacity can create an apparent raw diversity advantage without transport.	That every aspect of the learned maps is useless.
Matched-anchor rank favors GCN	The proposed common-frame representation-preservation advantage is not supported.	A causal mediation pathway from rank to prediction.
Random maps reduce BA by 8.86 points	Accepted predictions depend on properties of the learned map system.	Biological bundle geometry or causal neural transport.
Map shuffling has a smaller effect	The learned map collection matters more than the exact node assignment under this intervention.	That node identity is irrelevant in every architecture.
Degree-preserving rewiring reduces BA by 1.23 points	The constructed propagation graph contributes modestly to prediction.	Anatomical validity of the graph or a universal topology effect.
Diagnostic-performance correlations are weak	Large representation changes do not reliably co-vary with predictive changes across site-density cells.	Proof that representations can never mediate prediction.
E2 BuNN succeeds only in S1	Learned transport helps when topology dependence and fixed recoverable frames are deliberately present.	That ABIDE-I has the same planted geometry.
E2 controls fall to chance	Correct topology and recoverable frame information are necessary for this synthetic advantage.	A universal failure threshold for real connectomes.
Every S6 operator is perfect	Already-global inputs can remove the need for propagation in a controlled setting.	That ABIDE connectivity rows equal the synthetic summary.
Site and seed sensitivity	The small BuNN-GCN difference is not stable to all reasonable weighting and optimization choices.	That the full result is invalid or that site effects explain every difference.

The contrast against the connectome elastic net is more stable. Its interval remained below zero, indicating that the learned BuNN curve underperformed the regularized non-graph reference in this analysis. That is an important result because a paper comparing only neural graph variants could miss the strongest available baseline.

5.5 Limitations

The conclusions are conditional on ABIDE-I, C-PAC filtered data without global-signal regression, the AAL-116 atlas, Pearson correlations and Fisher transformation, positive-edge propagation graphs, full connectivity-row node features, the 0–20 percent density grid, the shared two-layer backbone, and the fixed tuning budget. Other atlases, preprocessing pipelines, graph construction rules, node features, or BuNN architectures may behave differently.

The study uses observational functional associations. It cannot identify anatomy, excitation, inhibition, or causal neural communication. The technical quality-control rule does not remove demographic or clinical sampling bias. Site sample sizes differ, and an equal-site mean answers a different question from a participant-weighted mean. The representation diagnostics are computational summaries; their association with prediction is not a causal mediation analysis.

E1 was specified only after the Study 1 result was known, so it is a post-hoc mechanistic audit rather than an independent confirmation. Its interventions are deliberately artificial. Random maps can disrupt the coordinate system expected by downstream learned weights, identity maps remove both learned orientation and reflections, and node shuffling preserves the learned map distribution while breaking its assignment. Their effect sizes reveal functional dependence under these checkpoints, not the effect of changing a single isolated biological mechanism. Finally, the completed report does not contain an external ABIDE-II validation cohort. A score-blind availability audit found the official 1,114-participant phenotype and quality records, but no complete C-PAC filtered, no-GSR AAL-116 ROI-time-series derivative for those participants. The alternative LLE repository documents mixed C-PAC and SPM-preprocessed volumes rather than the frozen Study 1 representation [13, 14]. The external evaluation gate therefore failed before any model was fitted or scored.

E2 was designed after the ABIDE-I result was known and cannot serve as an independent confirmation of that result. Its ring graph, frame marker, smooth signal, noise scale, and global summary are controlled computational objects, not estimates of brain organization. The near-perfect S0, S1, oracle, and S6 scores show that the intended positive controls were strong; they do not establish performance on less separable synthetic families. Learned BuNN also did not exactly recover the true relative transports, so E2 supports task-relevant alignment rather than

exact geometric identification.

The stopped v1 E2 run exposed an implementation error before unblinding: zero angles still produced reflections in half the bundles. The run was quarantined, an explicit identity regression test was added, and all v2 cells were trained from scratch. This correction improves the accepted evidence but also illustrates why generator-level tests are essential in mechanism studies.

5.6 Remaining extensions and their evidential role

The accepted-checkpoint intervention audit and synthetic known-geometry study are now complete and independently audited. Both remain separate from the confirmatory result because their protocols were frozen after Study 1 was known. Deeper matched GCN/BuNN networks and a leakage-free Wang-style time-series model remain separate studies. A prospective ABIDE-II evaluation was examined only at the metadata and derivative-availability stage. It remains blocked unless a matching official derivative appears or a separate raw-data preprocessing study is frozen and validated. E3 is only warranted if the paper needs to answer a remaining depth/capacity question; it is not needed to establish the conditional E2 mechanism.

6 Conclusion

In this pre-specified ABIDE-I analysis, learned orthogonal bundle transport did not improve the held-out-site performance–density curve over GCN. It also failed to surpass identity propagation or the learned-local control and remained below the regularized connectome baseline. Common-frame, matched-anchor diagnostics did not support the proposed representation-preservation advantage.

The finding does not show that bundle geometry is biologically false or that BuNNs are generally ineffective. It shows something narrower and useful: for dense functional-connectome graphs whose node features already contain global connectivity information, this learned transport operator added capacity and runtime without a detected advantage over GCN or elastic net. The post-hoc checkpoint audit adds an important qualification. Learned maps were computationally active: replacing them with random orthogonal maps reduced balanced accuracy by 8.86 points after multiplicity-controlled inference. Exact node-map placement and degree-preserving topology contributed less, and representation changes were only weakly associated with prediction changes.

The controlled E2 study supplies the missing positive result without changing the ABIDE-I conclusion. Learned BuNN gained 42.75 points on the pre-specified conditional contrast and

passed the oracle and learned-local gates. The advantage disappeared under wrong topology, shuffled frame information, unlearnable subject-specific frames, and already-global features. Bundle transport therefore helped when topology-dependent information had to be aligned through stable recoverable frames; it was unnecessary or ineffective when those conditions did not hold.

The combined evidence therefore explains the negative comparison without dismissing BuNN as inert or generally ineffective. It learned consequential internal transformations in ABIDE-I, and the same operator succeeded in a controlled task matching its intended boundary conditions. Those conditions did not yield better cross-site classification from the selected connectome graph and globally informed features. That distinction is the main contribution of the completed study. It was obtained through matched controls, held-out-site evaluation, strong non-graph baselines, gauge-aware diagnostics, frozen decision rules, checkpoint interventions, known-ground-truth controls, and independent artifact audits.

Evidence Boundary and Remaining Work

Table 11: Boundary between completed evidence and planned work.

Component	Completed and audited	Not yet a result
Data	ABIDE-I C-PAC filtered, no-GSR AAL time series; technical QC; 754-person connectome archive; score-blind ABIDE-II metadata and derivative-availability audit.	Alternative preprocessing, atlases, negative-edge propagation, and any ABIDE-II model evaluation.
Classical models	Covariate, connectome elastic-net, and combined baselines across all 18 held-out sites.	Additional harmonization or feature-selection families.
Neural models	Two-layer identity, learned-local, GCN, trivial-bundle, and learned-BuNN operator study; five seeds; complete density grid.	Five-layer matched study, Wang-inspired time-series architecture.
Mechanism evidence	Common-frame representation curves, matched anchors, site influence, seed stability, efficiency, the audited 360-checkpoint intervention study, and the 700-cell synthetic known-geometry study.	Optional matched depth/capacity sensitivity and component-level feature-loss analysis.
External validity	None claimed beyond the frozen ABIDE-I pipeline; the ABIDE-II compatibility gate failed before prediction.	Prospective ABIDE-II evaluation only after a matching derivative or separately frozen raw-data pipeline passes the gate.

A Reproducibility Record

The repository separates immutable scientific artifacts from presentation files. Raw participant derivatives, checkpoints, and credentials are excluded from the public release package. The report draws its numbers from generated tables and macros tied to accepted analysis hashes. The following checks were completed before the Study 1 result was interpreted:

- all 18 outer sites were present;
- expected tuning, prediction, metric, diagnostic, and runtime row counts were complete;
- model and data contracts matched their frozen hashes;
- hidden metrics were independently recomputed before unblinding;
- parallel worker outputs matched the validated sequential execution in the engineering equivalence test;
- site-bootstrap and sign-flip analyses were built and tested on synthetic data before use on accepted results;
- all 700 accepted E2 cells passed score-blind cell audits and all 100 condition-replicate groups were independently sealed before unblinding;
- the E2 analysis was invoked once under its frozen acknowledgment and independently recomputed artifact by artifact;
- report figures and tables were regenerated from the frozen result snapshot rather than edited manually.

The current source code, report source, manifests, and AI-use log are stored in the project repository. Participant-level source data are not embedded in this PDF.

B AI-Use Detail Required Before Submission

The local AI-use ledger is more detailed than the summary in the front matter, but a competition submission requires the researcher to verify it against retained chats and product records. At minimum, the final disclosure should state for every tool:

1. product and model/version where known;
2. dates or time period of use;

3. stages of the research in which it was used;
4. purposes and approximate frequency;
5. files, decisions, or outputs affected;
6. what the student independently checked; and
7. whether prior supervisor approval was obtained.

This report does not invent missing Claude or Cursor history. Those details must be supplied by the researcher.

References

- [1] K. Han, Y. Su, L. He, L. Zhan, S. Plis, V. Calhoun, and C. Yang, “Rethinking functional brain connectome analysis: Do graph deep learning models help?” *npj Artificial Intelligence*, vol. 2, p. 19, 2026. [Online]. Available: <https://doi.org/10.1038/s44387-025-00067-x>
- [2] J. Bamberger, F. Barbero, X. Dong, and M. M. Bronstein, “Bundle neural network for message diffusion on graphs,” in *International Conference on Learning Representations*, 2025. [Online]. Available: https://proceedings.iclr.cc/paper_files/paper/2025/hash/2ac879d1865475a7abc8dfc7a9c15c27-Abstract-Conference.html
- [3] C. Bodnar, F. Di Giovanni, B. P. Chamberlain, P. Li’o, and M. M. Bronstein, “Neural sheaf diffusion: A topological perspective on heterophily and oversmoothing in GNNs,” in *Advances in Neural Information Processing Systems*, vol. 35, 2022. [Online]. Available: https://proceedings.neurips.cc/paper_files/paper/2022/hash/75c45fca2aa416ada062b26cc4fb7641-Abstract.html
- [4] K. Murphy and M. D. Fox, “Towards a consensus regarding global signal regression for resting state functional connectivity MRI,” *NeuroImage*, vol. 154, pp. 169–173, 2017. [Online]. Available: <https://doi.org/10.1016/j.neuroimage.2016.11.052>
- [5] A. Di Martino *et al.*, “The autism brain imaging data exchange: Towards a large-scale evaluation of the intrinsic brain architecture in autism,” *Molecular Psychiatry*, vol. 19, no. 6, pp. 659–667, 2014. [Online]. Available: <https://doi.org/10.1038/mp.2013.78>
- [6] Preprocessed Connectomes Project, “ABIDE preprocessed,” INDI, Child Mind Institute, 2026, accessed 9 August 2026. [Online]. Available: https://fcon_1000.projects.nitrc.org/indi/abide/preprocessed.html
- [7] T. N. Kipf and M. Welling, “Semi-supervised classification with graph convolutional networks,” in *International Conference on Learning Representations*, 2017. [Online]. Available: <https://openreview.net/forum?id=SJU4ayYgl>
- [8] S. Parisot, S. I. Ktena, E. Ferrante, M. C. H. Lee, R. Guerrero, B. Glocker, and D. Rueckert, “Disease prediction using graph convolutional networks: Application to autism spectrum disorder and alzheimer’s disease,” *Medical Image Analysis*, vol. 48, pp. 117–130, 2018. [Online]. Available: <https://doi.org/10.1016/j.media.2018.06.001>
- [9] L. Wang, K. Li, and X. P. Hu, “Graph convolutional network for fMRI analysis based on connectivity neighborhood,” *Network Neuroscience*, vol. 5, no. 1, pp. 83–95, 2021. [Online]. Available: https://doi.org/10.1162/netn_a_00171
- [10] T. Yang, M. A. Al-Duailij, S. Bozdog, and F. Saeed, “Classification of autism spectrum disorder using rs-fMRI data and graph convolutional networks,” in *2022 IEEE International Conference on Big Data*, 2022, pp. 3131–3138. [Online]. Available: <https://doi.org/10.1109/BigData55660.2022.10021070>
- [11] N. Tzourio-Mazoyer, B. Landeau, D. Papathanassiou, F. Crivello, O. Etard, N. Delcroix, B. Mazoyer, and M. Joliot, “Automated anatomical labeling of activations in SPM using a macroscopic anatomical parcellation of the MNI MRI single-subject brain,” *NeuroImage*, vol. 15, no. 1, pp. 273–289, 2002. [Online]. Available: <https://doi.org/10.1006/nimg.2001.0978>
- [12] H. Zou and T. Hastie, “Regularization and variable selection via the elastic net,” *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, vol. 67, no. 2, pp. 301–320, 2005. [Online]. Available: <https://doi.org/10.1111/j.1467-9868.2005.00503.x>
- [13] Autism Brain Imaging Data Exchange, “ABIDE II release and collection documentation,” https://fcon_1000.projects.nitrc.org/indi/abide/abide_II.html, 2017, accessed 14 August 2026.
- [14] 1000 Functional Connectomes Project, “Locally linear embedding preprocessed repository,” https://fcon_1000.projects.nitrc.org/indi/abide/LLE_home.html, 2023, accessed 14 August 2026.